Toward Entity Retrieval over Structured and Text Data

Mayssam Sayyadian, Azadeh Shakery,
AnHai Doan, ChengXiang Zhai

Department of Computer Science

University of Illinois, Urbana-Champaign

Presentation at ACM SIGIR 2004 Workshop on Information Retrieval and Databases, July 29, 2004

Motivation

- Management of textual data and structured data is currently separated
- A user is often interested in finding information from both databases and text collections. E.g.,
  - **Course information may be stored in a database; course web sites are mostly in text**
  - **Product information may be stored in a database; product reviews are in text**
- How do we find information from databases and text collections in an integrative way?

Entity Retrieval (ER) over Structured and Text Data

- Problem Definition
  - **Given collections of structured and text data**
  - **Given some known information about a real-world entity**
  - **Find more information about the entity**
- Example
  - **Data= DBLP (bib. Database) + Web (text)**
  - **Entity = researcher**
  - **Known information = "name of researcher" and/or a paper published by the researcher**
  - **Goal = find all papers in DBLP and all web pages mentioning this researcher**

Entity Retrieval vs. Traditional Retrieval

- ER vs. Database Search
  - **ER requires semantic-level matching**
  - **DB search matches information at the syntactic-level**
- ER vs. Text Search
  - **ER represents a special category of information need, which is more objectively defined**
- What's new about ER?

Challenges in ER

- Requires semantic-level matching
  - **Both DB search and text search generally match at the syntactic level**
  - **E.g., name= "John Smith" would return all records match the name in DB search**
  - **E.g., query="John Smith" would return documents match one or both words**
  - **But "John Smith" could refer to multiple real-world entities**
- Same name for different entities
- A unique entity name may appear in different syntactic forms in a DB and text collection.
  - **E.g., "John Smith" -> "J. Smith"**

Definition of a Simplified ER Problem

Query $Q=(q, R, C, T)$

q=Text query

$R=\{r_1, r_2, \ldots, r_m\}$ examples of rel docs

$r_i \in D$

$C=\{c_1=v_1, c_2=v_2, \ldots, c_n=v_n\}$ constraints

$c_i \in A$

$T=\{t_1, t_2, \ldots, t_l\}$ target attributes

$t_i \in A$

Data

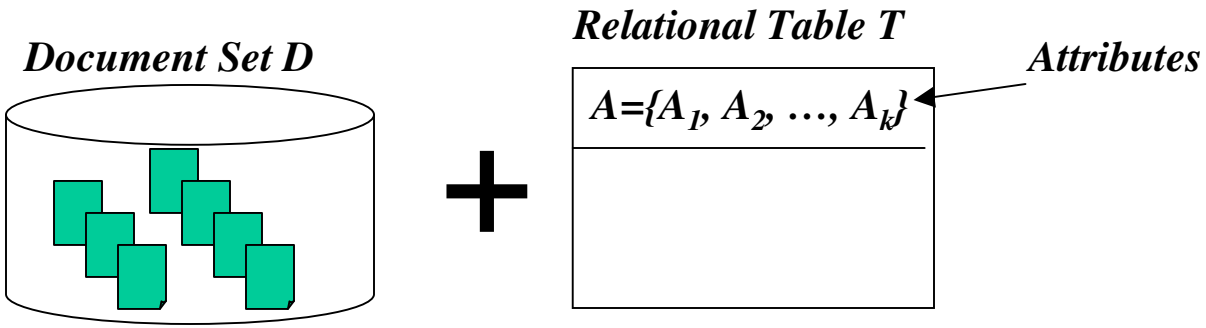

Results

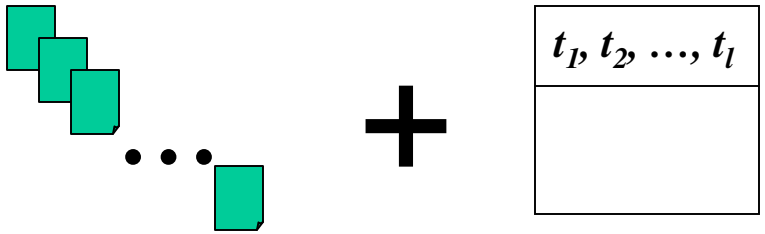

Finding all Information about "John Smith"

Query *Q=(q, R, C, T)*

q="John Smith"

R: Home page of "John Smith"

C: {author="John Smith", paper.conferenc=SIGIR}

T: {paper.title, paper.conference}

Data

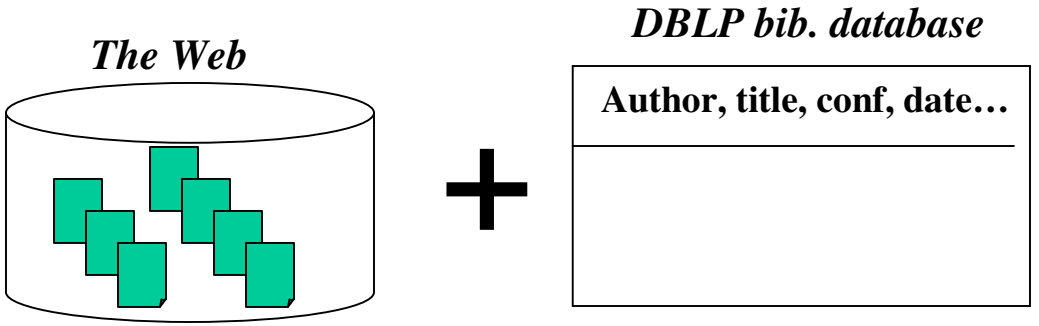

"John Smith" is highly ambiguous!

Results

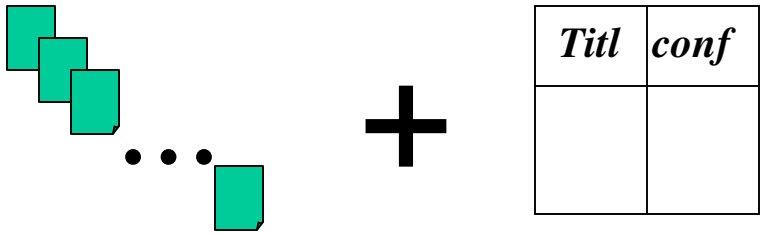

ER Strategies

- Separate ER on DB and on text
  - **Q=(q,R,C,T)**
    - Use Q1=(q,R) to search the text collection
    - Use Q2=(C,T) to search the DB
  - **The main challenge is entity disambiguation**
- Integrative ER on DB + Text
  - **Q=(q,R,C,T): use Q to search both the text collection and DB**
  - **Relevant information in DB can help improve search over text**
  - **Relevant information in text can help improve search over DB**

Hypothesis tested in this work

Exploit Structured Information to Improve ER on Text

Given an ER query Q=(q,R,C,T)
Assume that we have a basic text search engine
We may exploit structured information to construct a different text query Q_i

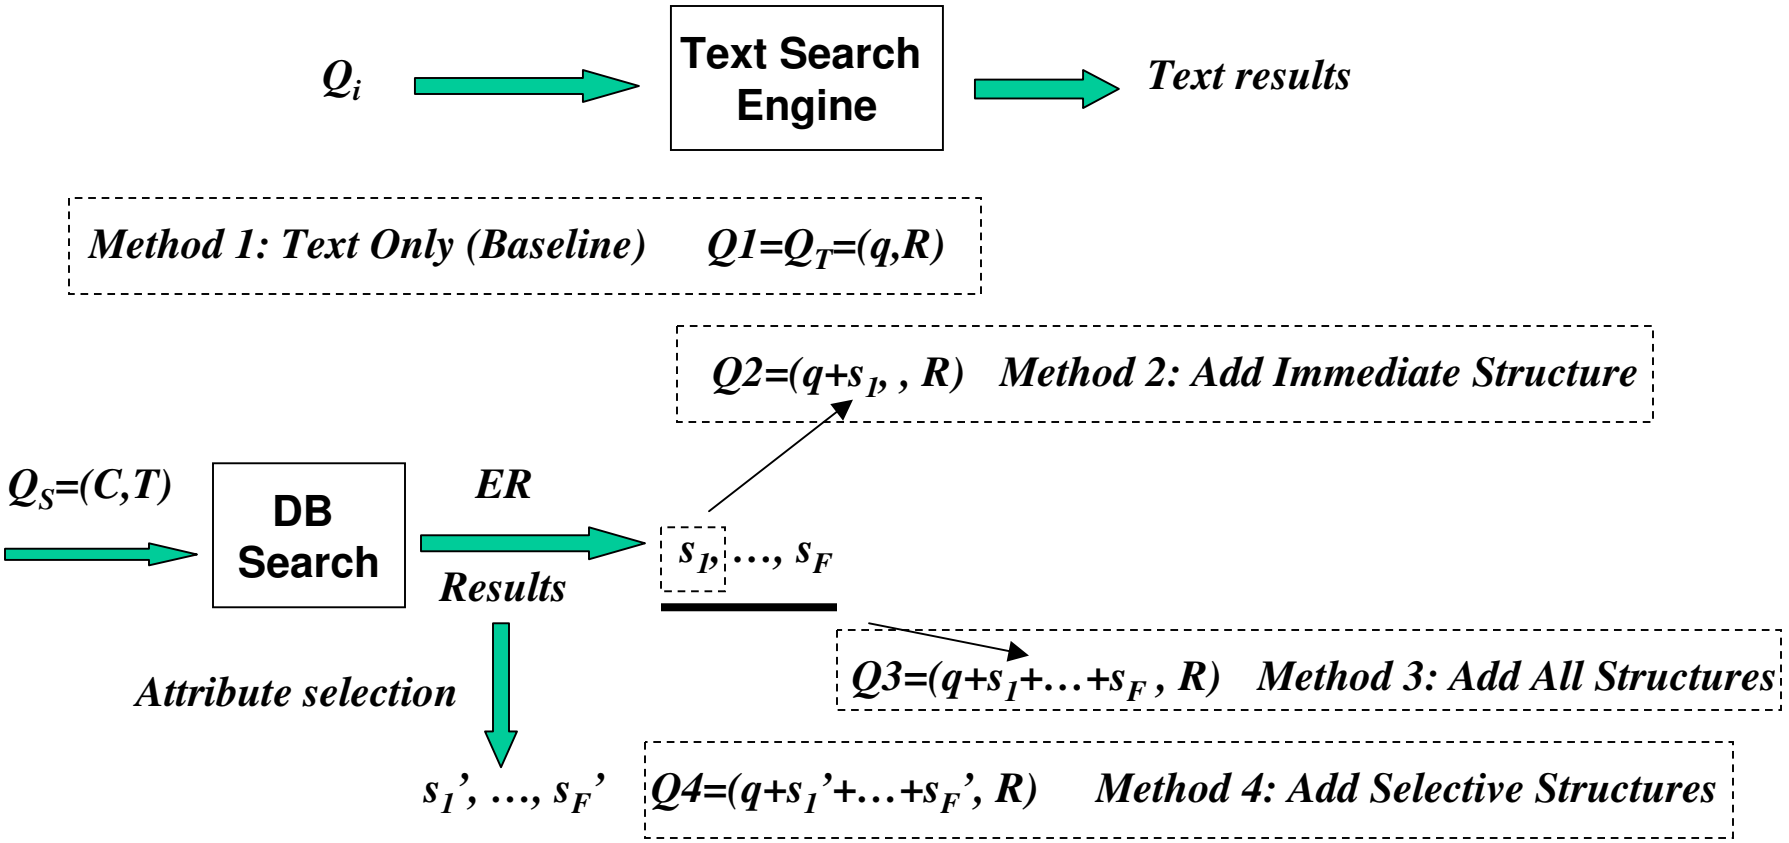

Attribute Selection Method

- Assumption: An attribute is more useful if it occurs more frequently in the top text documents (returned by the baseline TextOnly method)
- Attribute Selection Procedure
  - **Use the top 25% of the docs returned by TextOnly as the reference doc set**
  - **Score each attribute by the average frequency of all the attribute values of the attribute in the reference doc set**
  - **Select the attribute with the highest score to expand the query**

Experiments

- ER queries: 11 researchers, Q=name (no relevant text doc examples)
- DB = DBLP (www.informatik.uni-trier.de/ley/db) , >460,000 articles
- Text collection = top 100 web pages returned by Google using the names of the 11 researchers
- Measures:
  - **Precision: percent of pages retrieved that are relevant**
  - **Recall: percent of relevant pages that are retrieved**
  - **F1: a combination of precision and recall**
- Retrieval method
  - **Vector space model with BM25 TF**
  - **Scores normalized by the score of the top-ranked document**
  - **A score threshold is used to retrieve a subset of the top 100 pages returned by Google (set to a constant all the time)**
  - **Implemented in Lemur**
- ER on DB: the DBLP search engine on the Web with manual selection of relevant tuples

Effect of Exploiting Structured Information

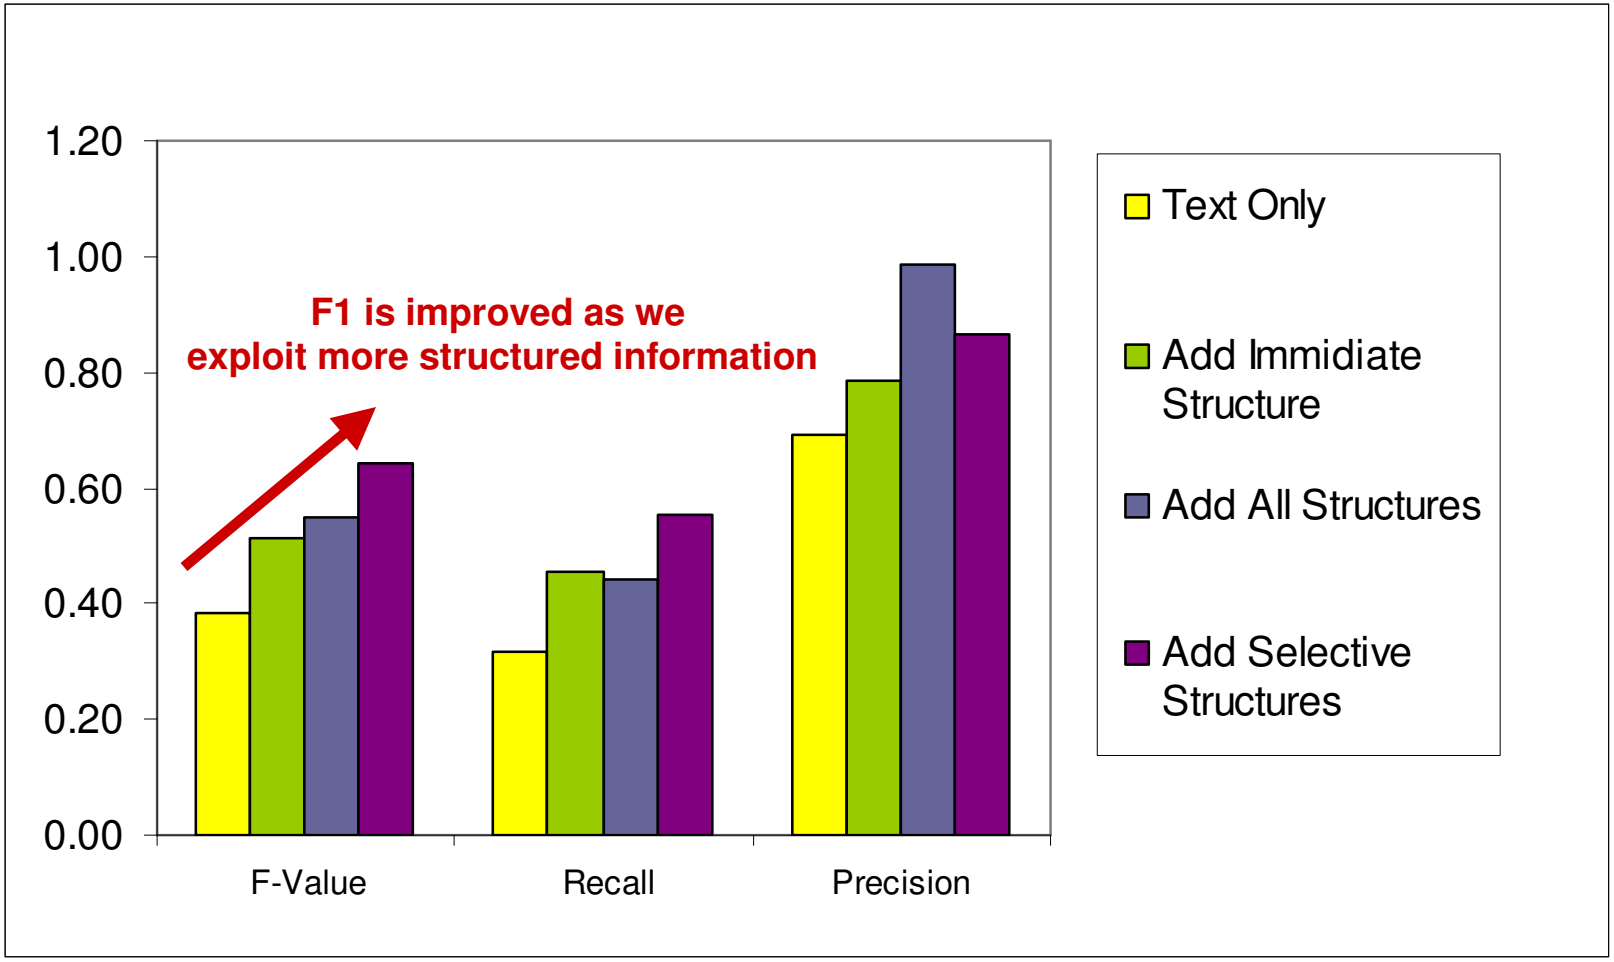

Effect of Attribute Selection

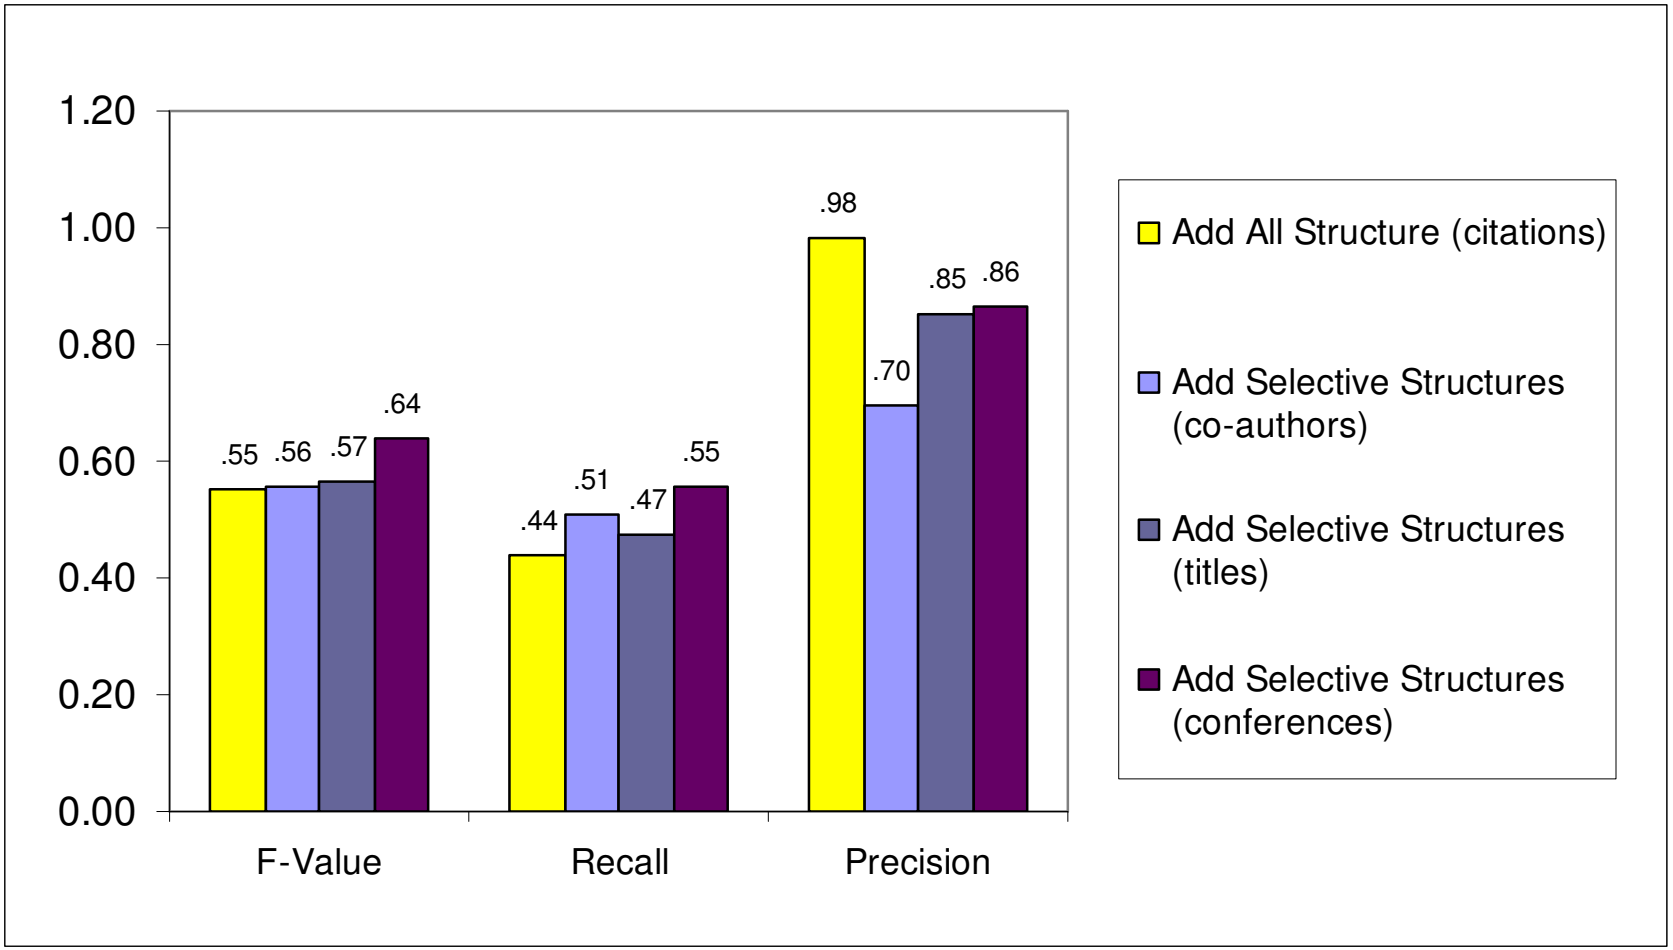

Conference is a better attribute than co-authors or titles

Automatic Attribute Selection

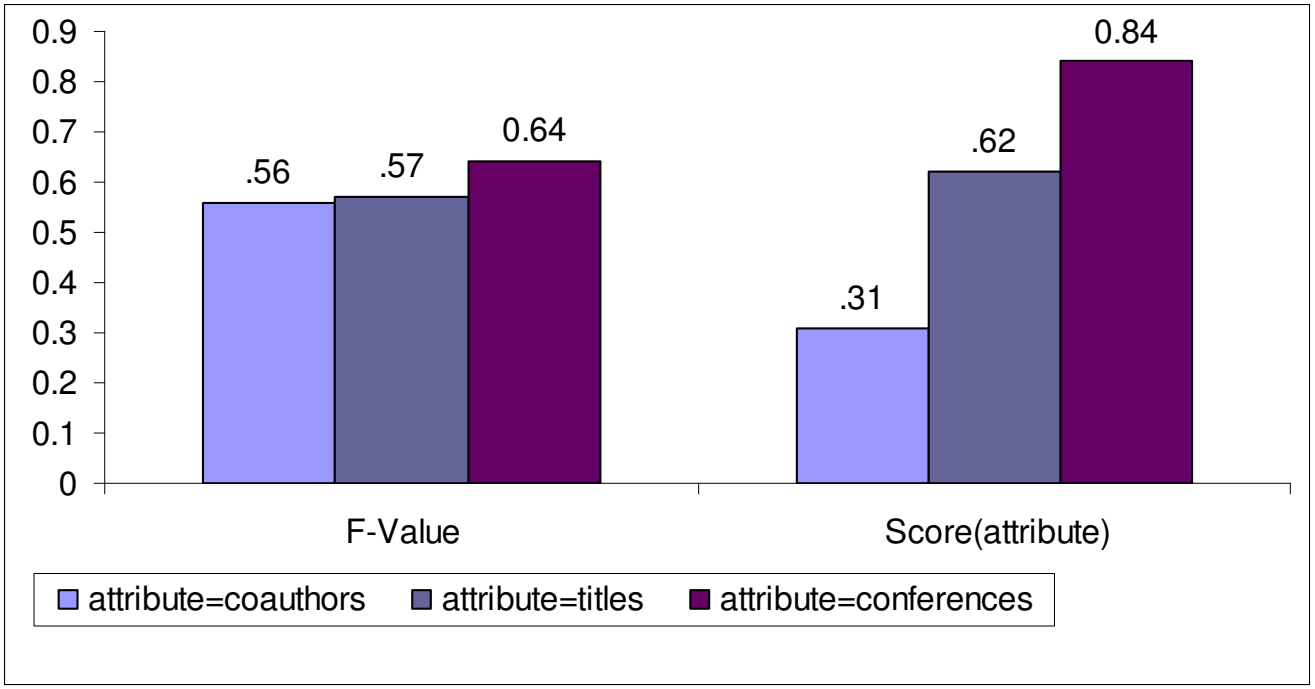

The attribute score based on value frequency predicts the usefulness of an attribute well

Conclusions

- We address the problem of finding information from databases and text collections in an integrative way
- We introduced the entity retrieval problem and proposed several methods to exploit structured information to improve ER on text
- With some preliminary experiment results, we show that exploiting relevant structured information can improve ER performance on text

Many Further Research Questions

- What is an appropriate query language for ER?
- What is an appropriate formal retrieval framework for ER?
- What are the best strategies and methods for ER?
- …

Thank You!